VERA — Vision-Enabled Reasoning Agent — System Architecture

v1.70.2 — Calibrated fresh-frame visual servoing

Root-owned deployment controller (separate from VERA)
vera update + authenticated System GUI /run/vera-updater/control.sock peer-checked fixed operations
vera-updater.timer vera-updater-check.service local-time installation window
local Unix socket
vera-updater.service /opt/vera-updater/src/updater stage / deploy / verify / recover
manifest + verified assets
atomic journal
ZeroMQ IPC
verified full-duty kernel PWM before thermal restart
GitHub stable releases mazum01/desktop-assistant HTTPS, exact tag/commit
/var/lib/vera-updater protected backups + transaction state post-reboot continuation
deployment.readiness RPC loaded version/revision/PID fresh telemetry + service/hardware health
NF-A6x25 PWM Fan PWM ← GPIO13 (kernel HW PWM, 25 kHz) TACH → GPIO6 (2 ppr)
tach
PWM
Hardware (Raspberry Pi 5)
Waveshare ESP32-C6-LCD-1.47 172×320 ST7789 SPI TFT (BLE mouth/status display + BLE OTA)
reSpeaker XVF3800 output → attached speaker
Pi Cam Module 3 Wide (CSI-0)
reSpeaker XVF3800 (USB mic array)
Hailo AI Accelerator (Hailo-8 / Hailo-10H, PCIe)
TMP117 (VERA3, I²C 0x48)
Pan servo + optional tilt PCA9685 configurable channels 180deg direct-drive / legacy DS3218 configured safe travel limits
TMP102 (VERA, I²C 0x48)
CSI
USB
PCIe
I²C
I²C/PWM
I²C
Drivers (src/)
vision/camera.py (picamera2)
audio/input.py
perception/hailo_models.py config/hailo/models.yaml probe / select / download HEFs python3 -m ...hailo_models --install
audio/version_announcer.py
perception/hailo_inference.py (VStreams + InferModel API)
thermal/tmp117.py
motion/visual_servo.py stable target lock + damped visual feedback identified camera/servo Jacobian config/tracking_calibration.json (local)
motion/servo_controller.py
USB
matched SCRFD HEF
TMP117 ID
perception/face_detector.py (SCRFD decode + Haar fallback)
audio/tts.py (espeak-ng)
perception/object_detector.py (YOLOv11l on-chip NMS; YOLO26m+ONNX H8 fallback)
thermal/temperature_sensor.py (probes TMP117 / TMP102)
thermal/fan_tach.py (lgpio edge cb)
thermal/fan.py (sysfs HW PWM + lgpio fallback)
renders to
audio/output.py
thermal/thermal_manager.py
SPI + GPIO
matched depth HEF
matched ArcFace HEF
External processes
firmware/upload_esp32_display_ble.sh + nimbleota.py uploader (BLE firmware push, no USB)
Web Browser http://<pi>:8080
~/.local/share/desktop-assistant/ telemetry.db (SQLite ring) faces.db (identity registry) thumbs/ (64×64 JPEG face crops)
scripts/desktop-assistant (CLI: ping / status / pan / say / watch / ...)
3rd-party subscribers (dashboards, future ML pipelines)
HTTP/WS :8080
latest_frame()
writes telemetry.db
latest_chunk() + chunk_index()
ZMQ PUB→SUB (cross-process)
ZMQ PUB→SUB (cross-process)
writes faces.db
ZMQ REQ/SUB
latest_frame()
ZMQ SUB
BLE GATT (OTA)
VisionService → vision.frame_ready ← vision.capture_still + latest_frame()
latest_frame()
WebService (FastAPI) GET / — dashboard SPA GET /stream — MJPEG 640×360 feed WS /ws — live JSON status GET /api/faces — list (has_thumb flag) GET /api/skills — registered skills GET /api/faces/{id}/thumb — face JPEG PUT /api/faces/{id} — rename DEL /api/faces/{id} — delete one DEL /api/faces — delete all GET|PUT /api/settings/quiet-hours GET|PUT /api/settings/audio GET|PUT /api/settings/voice GET|PUT /api/room GET|PUT /api/music/* — proxied via IPCClient GET|POST /api/podcasts/* — proxied via IPCClient POST /api/say|pan|version|utterance port 8080
IPCClient RPC (music/podcast reads+actions)
IPCBridge (media) PUB ipc:///tmp/desktop-assistant-media.pub REP ipc:///tmp/desktop-assistant-media.rep SUB ← core.pub (upstream) RPCs: music.get_state/set_volume/... podcast.search/subscribe/play/...
DisplayService ← system.startup_status / service.started / *.error ← display.spectrum / av.speaking_started / av.spoke → display.status / display.error (BLE write to ESP32: status / mouth / eq)
PlaybackSpectrumService ← music.state_changed / podcast.playback pw-record on sink monitor + FFT → display.spectrum
FaceService ← perception.faces / face.meet → av.say (greetings) → face.identified cooldown-based re-greet
AudioCaptureService → audio.level / chunk + latest_chunk()
TelemetryService ← thermal.* / motion.position / audio.level → telemetry.flush SQLite ring buffer
PerceptionService ← vision.frame_ready → perception.faces / perception.error + FaceEmbedder (ArcFace) + FaceRegistry
VoiceCommandService ← audio stream/VAD → voice.state / wake / transcript / intent → av.utterance (routed transcript) wake-window + dialog context
systemd: desktop-assistant-thermal.service
IPCBridge (thermal) PUB ipc:///tmp/desktop-assistant-thermal.pub REP ipc:///tmp/desktop-assistant-thermal.rep
PodcastService Apple Podcasts search + RSS subscriptions ← podcast.play/pause/resume/stop/refresh → podcast.playback / subscriptions_updated
MusicService ← music.play / stop / pause / next ← music.thumbs_up / thumbs_down / set_volume Pandora / audio playback
ClockService + ClockAnnouncer → av.say (hourly/half-hourly) respects QuietHours
NotificationService ← thermal.temp / perception.faces ← av.speaking_started / av.spoke → av.say (thermal alerts, absence alerts) rate-limited · quiet-hours aware
SkillsService ← av.utterance / thermal.temp → av.say / vision.describe / av.tell_joke → face.meet / motion.pan_to / music.* → object.set_enabled / music.set_volume → tracking.set_face_tracking SkillRegistry — first-match dispatch 17 skills: help · system_status · volume quiet_hours · weather · reminder smart_home · news + 9 others ConfigField framework: enabled / config_schema
IPCBridge (core) PUB ipc:///tmp/desktop-assistant.pub REP ipc:///tmp/desktop-assistant.rep SUB ← thermal.pub (upstream) cmds: ping / topics / last / publish / status
AVService ← av.say / beep / utterance / announce_version → av.spoke / av.speaking_started / version_announced
ThermalService — thermal.temp — thermal.fan — thermal.rpm — thermal.critical — thermal.error
TrackingService ← perception.faces (atomic frame time/geometry) ← motion.position / settings / limits_changed → motion.pan_to / tilt_to (optional) ← tracking.start_autotune / cancel_autotune → motion.calibration_lock / av.say ← av.spoke (announce before probes) → tracking.debug / autotune_progress / autotune_done fresh observations only, settle then correct Tilt: travel limits and center +/-20deg idle random pan · speaking nod
MotionService shared PCA9685, independent pan / tilt 50Hz bounded quintic profiles ← motion.pan_to / tilt_to / center / relax / stop ← motion.calibration_lock (expiring lease) ← motion.set_channels / set_limits / set_tilt_limits → motion.position / moved / settings / error → motion.channels_changed / limits_changed / tilt_limits_changed RPC motion.get_status / motion.set_channels
ObjectService ← vision.frame_ready / object.set_enabled → object.detections 2fps · max_objects cap · BGR→RGB fix
latest_frame()
vision.mono_depth_map
latest_frame()
BLE GATT status / mouth / eq
config R/W
suppresses greetings
suppresses announcements
suppresses alerts
re-emit
extends
systemd: desktop-assistant-core.service
MonoDepthService scdepthv3 on Hailo-8 / Hailo-10H → vision.mono_depth_map relative 320×256 depth, 3 Hz
Arduino firmware firmware/vera_display (ESP32-C6 process boundary) NimBLEOta service (BLE OTA, 0x8018)
RoomService ← room.set → room.updated / av.say persists room name · visual scene baseline histogram · sustained divergence detection → TTS prompt
Core (src/core/)
version.py (reads /VERSION)
media MessageBus (in-process pub/sub)
QuietHours (enabled / start / end loaded from config/quiet_hours.yaml)
Service base (start/tick/stop)
core MessageBus (in-process pub/sub)
thermal MessageBus (in-process pub/sub)